

# Auto-Complete Query Retrieval System

Team: **No Direction**

Dice Challenge 2.0

# The Challenge: Real-Time Auto-Complete at Scale

## Goal

Suggest the **top 150** most likely query completions for any user prefix.

## Dataset & Constraints:

- 522,726 test prefixes
- 4.2M candidate queries
- 5.3M query features
- Time limit: **90 minutes**
- Memory: 32GB RAM
- CPU: 32 cores available

**Evaluation:** Hit@150 target: **70–80%** (competition goal).

# Initial Attempt: Only BM25 (Character Trigrams)

- Tokenize queries into character n-grams ( $n=3$ ) to handle typos.
- Build BM25 index on 4.2M queries and retrieve top 150 candidates per prefix.

## Catastrophe: 24 DAYS Runtime

- Processing speed: **0.25 prefixes/second**
- Total time needed:  $522,726 \div 0.25 = 2,090,904$  seconds **24.2 days!**
- Root causes: single-threaded pipeline, large candidate pool, validation overhead, execution on Laptop CPU (no parallelization)

# Character Trigrams — Why We Used Them

- Handle typos and partial words more gracefully than word tokens.
- Example (partial word): smar → trigrams [sma, mar] match smartphone.
- Example (typo): blak → overlaps with black via bla.

**Used as the lexical backbone for BM25 candidate generation; later complemented by prefix-aware heuristics and fuzzy matching.**

# Optimization 1: Parallelization + Aggressive Sampling

- Used 32 CPU cores (`multiprocessing.Pool`) using SSH.
- Reduced candidate pool: 4.2M  $\rightarrow$  50K (random sampling).
- Disabled expensive validation during runs to prioritize throughput.

## Result: 72 Minutes (1600x Speedup)

- Processing speed: **138 prefixes/second**
- Total time: **72 minutes**
- Output: `submission.csv` (1.5GB)

# The Sampling Dilemma

**Pitfall: Very less matched in many cases.**

Random 50K sampling produced prefixes with **zero** relevant candidates (e.g., “sopt ha”).

## Analysis:

- 50K sample = **1.2%** coverage of 4.2M queries
- Random sampling doesn't guarantee inclusion of rarer but relevant queries
- Result: Poor Results despite fast runtime

# Problem: Speed Without Quality

## Quality Crisis

Results from the 50K sample were largely **irrelevant** for many prefixes.

### Examples of poor matches (50K pool):

Prefix	Pure BM25 Result
"sopt ha"	"cotton saree pink sopt"
"harf pent"	"q pental penpencil 07"
"pink partywear s"	"partywear sofa"
"black farshi salwar"	"printed farshi salwar suit"

**Root cause:** BM25 matches trigrams anywhere; no prioritization of prefix-starting candidates.

# Memory Crisis: Full 4.2M Pool Attempt

- BM25 index size (4.2M): ~4.5GB
- 32 workers  $\times$  4.5GB = **144GB** required
- System crashed / thrashing; process killed by OS

## Conclusion

Full pool per worker is infeasible under memory constraints (32GB total).



# Solution: Bucketing + Balanced Sampling

## Bucketing by Match Type

Categorize candidates first, then rank within each bucket.

Three priority buckets:

- 1 **Exact Matches:** `candidate.startswith(prefix)`
- 2 **Contains Matches:** `prefix in candidate`
- 3 **BM25 Alternatives:** lexical similarity / typos

Top 150 composition used:

$$\text{Top 150} = \text{exact}[: 100] + \text{contains}[: 40] + \text{others}[: 10]$$

## Balanced Pool: 500K Sweet Spot

- Increased candidate pool to **500K** (10x of 50K)
- Memory per worker: ~200MB (manageable)
- Coverage: **12%** of full pool (vs 1.2%)
- Expected quality improvements: **40–80/150 exact matches (varies by prefix)**

**Resulting runtime: 65 minutes**, stable (within 90 minute limit).

# Performance Evolution

Approach	Time	Speed	Quality / Status
Pure BM25 (initial)	Too Long	0.25/s	Failed
+ Parallel (50K pool)	72 min	138/s	Low (poor quality)
+ Smart reranking	83 min	138/s	Medium (better)
+ Full pool (4.2M)	—	—	OOM crash
+ <b>500K pool (final)</b>	<b>65 min</b>	<b>134/s</b>	<b>Good (Success)</b>

# Quality Comparison: 50K vs 500K Pool

Prefix	50K Pool (Top-5 exact)	500K Pool (Top-5 exact)
“black farshi salwar”	1/5	4/5
“sopt ha”	0/5	2/5
“pink partywear s”	0/5	4/5
“dress for”	2/5	5/5
Average Top-5	<b>20%</b>	<b>80%</b>
Average Top-150	<b>30%</b>	<b>60–75%</b>

# Sample Output: Before vs. After

## Before (50K pool):

```
1 Prefix: 'black farshi salwar'
2 Top 5:
3   [ ] printed farshi salwar
4   [ ] parsi salwar set
5   [ ] chunni salwar colour
6   [ ] karachi suit salwar
7   [ ] toy wali salwar
8
9 0/5 exact | 0/150 total
10
```

## After (500K pool):

```
1 Prefix: 'black farshi salwar'
2 Top 5:
3   [ ] black farshi salwar suit
4   [ ] black farshi salwar kameez
5   [ ] black farshi salwar set
6   [~] printed black farshi salwar
7   [ ] dark farshi salwar
8
9 3/5 exact | 65/150 total
10
```

## Main Idea: Combine multiple ways of thinking about similarity

- **BM25 (Text Match):** Looks at overlapping words or characters — good for exact matches and prefixes.
- **Semantic Model:** Uses deep learning embeddings (Sentence Transformers) to find meaning-based matches, even if words differ.
- **Popularity:** Gives higher priority to queries people click or buy from more often.
- **Memorization:** If a prefix has been seen before in training data, boost those known results.

## How it works

- 1 Build two models: a fast text-based (BM25) and a meaning-based (semantic) one (all-MiniLM-L6-v2)
- 2 For each prefix:
  - Find matching or similar queries.
  - Score them by combining text, meaning, and popularity.
  - Boost known good ones and keep the top few suggestions.
- 3 Save the ranked results for use in an autocomplete system.

But, 1.3sec/iteration, which is impractical for real world usage.

# Top 3 Candidate Queries

Prefix: full\_sleeve\_mehendi\_st

- 1 full sleeve mehendi stick
- 2 full sleeve mehendi stencil
- 3 full sleeve mehendi sticker

Prefix: vridavan\_dress

- 1 vridavan dress
- 2 vridavan dress for girl
- 3 vrindavan dress

Prefix: lucifer\_w

- 1 lucifer watch
- 2 lucifer watch combo
- 3 lucifer watch chain

Prefix: parryware

- 1 parryware toilet
- 2 parryware tap cleaner
- 3 parryware toilet seat cover

Prefix: suzume

- 1 suzume
- 2 suzume novel
- 3 suzume 3 manga

Prefix: heavy\_readymad

- 1 heavy readymade suit
- 2 heavy readymade blouse
- 3 heavy readymade salwar suit

Prefix: orange\_color\_su

- 1 orange color suit
- 2 orange color suits
- 3 orange color suit pant

Prefix: mars\_moisturize

- 1 mars moisturizer
- 2 moisturize
- 3 moisturize oil



# New Approach: A Two-Stage "Filter & Rank" Pipeline

- **Stage 1: Candidate Generation (Recall-Focused)**

- Cast a wide net to retrieve ~500 potentially relevant queries.
- Goal: Ensure the correct query is captured in this initial set.
- Methods: Hybrid approach using Lexical and Semantic search.

# New Approach: A Two-Stage "Filter & Rank" Pipeline

- **Stage 1: Candidate Generation (Recall-Focused)**

- Cast a wide net to retrieve ~500 potentially relevant queries.
- Goal: Ensure the correct query is captured in this initial set.
- Methods: Hybrid approach using Lexical and Semantic search.

- **Stage 2: Re-Ranking (Precision-Focused)**

- Intelligently score and sort the ~500 candidates.
- Goal: Push the single best query to the top of the final list of 150.
- Method: A powerful LightGBM machine learning model.

# Ideal Solution Pipeline




Figure: The complete two-stage pipeline from prefix to final ranked list.

# Stage 1 Deep Dive: Hybrid Candidate Generation

## Lexical Search: The Typo Catcher

**Method:** BM25 with Character N-grams (we used tri-grams)

**Why?:** By breaking words into small character chunks (e.g., "shoes" → ['sho', 'hoe', 'oes']), this method finds matches even with severe typos.

*Example: A typo like "shos" will still have a high score for the query "shoes".*

## Semantic Search: The Meaning Matcher

Sentence-BERT + FAISS

**Why?:** This understands the user's intent, not just the letters they typed. It finds queries with similar meanings, even if they use different words.

*Example: A prefix like "clothing for ladies" will find the query "dress for women".*

# Improvement: Smart Reranking with Prefix Bonus

## Insight

Autocomplete should prioritize **prefix matches** over arbitrary substring matches.

Hybrid scoring (applied to BM25 candidates):

```
1 for candidate, bm25_score in bm25_candidates:
2     if candidate.startswith(prefix):
3         final_score = bm25_score * 2.0
4     elif prefix in candidate:
5         final_score = bm25_score * 1.5
6     else:
7         final_score = bm25_score * 0.8
8
```

## Stage 2 Deep Dive: Precision Re-Ranking

### The "Expert Judge": LightGBM Model

We trained a LightGBM model to act as an expert judge. It analyzes each of the ~500 candidates and assigns a precise relevance score.

### Key Features (The Evidence)

The model's decisions are based on a rich set of features:

- **Semantic Similarity:** The cosine similarity between prefix and candidate embeddings.
- **Lexical Score:** The relevance score from our BM25 index.
- **Popularity Metrics:** Historical performance data like orders, clicks, and volume from the provided datasets.
- **String Features:** Simple but effective metrics like length ratios and character overlap.

# Final Approach: Multi-Strategy, Prefix-First Retrieval

## Core Idea

Prioritize reliable prefix candidates, then expand with fuzzy/lexical fallbacks.

### Retrieval strategies (in order):

- ➊ **Historical matches** (highest weight): learn which queries users actually searched for given a prefix historically.
- ➋ **Direct prefix index**: fast lookup for queries starting with prefix (trie-like behavior).
- ➌ **Shorter-prefix expansion**: try shorter prefixes to improve recall.
- ➍ **Fuzzy matches**: lightweight fuzzy matching (e.g., rapidfuzz) for typos.
- ➎ **Popular queries fallback**: global popular queries if above yield few results.

# Results Summary

- **Runtime: 65 minutes** (within 90 min limit)
- **Throughput: 134 prefixes/second**
- **Memory: Stable** (no crashes)
- **Output: `submission_fast.csv`** (1.5GB)
- **Speedup from baseline: 531x** (24 days → 65 minutes)